

# Train model

From cs231n

August 2024

## Things you might try

To train a ConvNet model, you can experiment with architectures, hyperparameters, loss functions, and optimizers to train a model.

### Things you might try:

- **Filter size:** Above we used 5x5; would smaller filters be more efficient?
- **Number of filters:** Above we used 32 filters. Do more or fewer do better?
- **Pooling vs Strided Convolution:** Do you use max pooling or just stride convolutions?
- **Batch normalization:** Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- **Network architecture:** The network above has two layers of trainable parameters. Can you do better with a deep network? Good architectures to try include:
  - conv-relu-pool  $\times N \rightarrow$  affine  $\times M \rightarrow$  softmax or SVM
  - conv-relu-conv-relu-pool  $\times N \rightarrow$  affine  $\times M \rightarrow$  softmax or SVM
  - batchnorm-relu-conv  $\times N \rightarrow$  affine  $\times M \rightarrow$  softmax or SVM
- **Global Average Pooling:** Instead of flattening and then having multiple affine layers, perform convolutions until your image gets small (7x7 or so) and then perform an average pooling operation to get to a 1x1 image picture (1, 1, Filter#), which is then reshaped into a (Filter#) vector. This is used in <https://arxiv.org/abs/1512.00567> Google's Inception Network (See Table 1 for their architecture).
- **Regularization:** Add l2 weight regularization, or perhaps use Dropout.

## Tips for Training

For each network architecture that you try, you should tune the learning rate and other hyperparameters. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations.
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.
- You should use the validation set for hyperparameter search, and save your test set for evaluating your architecture on the best parameters as selected by the validation set.

## Going Above and Beyond

If you are feeling adventurous there are many other features you can implement to try and improve your performance. You are **not required** to implement any of these, but don't miss the fun if you have time!

- Alternative optimizers: you can try Adam, Adagrad, RMSprop, etc.
- Alternative activation functions such as leaky ReLU, parametric ReLU, ELU, or MaxOut.
- Model ensembles.
- Data augmentation.
- New Architectures:
  - <https://arxiv.org/abs/1512.03385> ResNets where the input from the previous layer is added to the output.
  - <https://arxiv.org/abs/1608.06993> DenseNets where inputs into previous layers are concatenated together.
  - <https://chatbotlife.com/resnets-highwaynets-and-densenets-oh-my-9bb15918ee32> This blog has an in-depth overview.

**Have fun and happy training!**